

NAME : RETHINAGIRI G

ROLL NO : 225229130

COURSE TITLE : PRACTICAL MACHINE LEARNING LAB

LAB-05 . Diabetes Classification using Logistic Regression

Step1. Import Data

```
In [1]: import pandas as pan
```

```
In [2]: dbs=pan.read_csv("C:\\Users\\user\\Downloads\\dataset_pml\\diabetes.csv")
```

```
In [3]: dbs.head()
```

Out[3]:

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
0	6	148	72	35	0	33.6	0.627	50	1
1	1	85	66	29	0	26.6	0.351	31	0
2	8	183	64	0	0	23.3	0.672	32	1
3	1	89	66	23	94	28.1	0.167	21	0
4	0	137	40	35	168	43.1	2.288	33	1

```
In [4]: dbs.shape
```

Out[4]: (768, 9)

```
In [5]: dbs.columns
```

Out[5]: Index(['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness', 'Insulin', 'BMI', 'DiabetesPedigreeFunction', 'Age', 'Outcome'], dtype='object')

```
In [6]: type(dbs)
```

Out[6]: pandas.core.frame.DataFrame

```
In [7]: dbs.dtypes
```

Out[7]: Pregnancies int64
Glucose int64
BloodPressure int64
SkinThickness int64
Insulin int64
BMI float64
DiabetesPedigreeFunction float64
Age int64
Outcome int64
dtype: object

```
In [8]: dbs.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 768 entries, 0 to 767  
Data columns (total 9 columns):  
#   Column                Non-Null Count  Dtype  
---  ---  
0   Pregnancies           768 non-null   int64  
1   Glucose                768 non-null   int64  
2   BloodPressure          768 non-null   int64  
3   SkinThickness          768 non-null   int64  
4   Insulin                768 non-null   int64  
5   BMI                    768 non-null   float64  
6   DiabetesPedigreeFunction 768 non-null   float64  
7   Age                    768 non-null   int64  
8   Outcome                768 non-null   int64  
dtypes: float64(2), int64(7)  
memory usage: 54.1 KB
```

```
Out[9]: <bound method IndexOpsMixin.value_counts of 0
1      0
2      0
3     94
4    168
...
763   180
764    0
765   112
766    0
767    0
Name: Insulin, Length: 768, dtype: int64>
```

```
In [10]: X=dbs.drop(["Outcome"],axis=1)
          X
```

768 rows \times 8 columns

```
Out[11]: array([[1, 0, 1, 0, 1, 0, 1, 0, 1, 1, 0, 1, 0, 1, 1, 1, 1, 0, 1, 0, 0,
1, 1, 1, 1, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 1, 1, 0, 0, 0, 1,
0, 1, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 1, 0,
1, 0, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0,
1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1,
1, 1, 0, 0, 1, 1, 1, 0, 0, 0, 1, 0, 0, 0, 1, 1, 0, 0, 1, 1, 1, 1,
1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0,
1, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 1, 0, 0, 0, 1, 1, 1, 0, 1, 1,
1, 1, 0, 0, 0, 1, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 1, 1, 0, 0, 0, 1,
0, 1, 0, 1, 0, 0, 0, 0, 0, 1, 1, 1, 1, 0, 0, 1, 1, 0, 1, 0, 1,
1, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 1, 0, 0, 0, 1, 1, 1, 1, 0, 0,
1, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 1, 0, 1, 0, 0,
1, 0, 1, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 1, 1,
0, 1, 0, 0, 1, 1, 1, 0, 0, 1, 0, 1, 0, 1, 1, 0, 1, 0, 1, 0, 0, 1,
1, 1, 0, 0, 1, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0,
0, 0, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0,
0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 1, 0, 0, 1, 0, 0,
0, 1, 0, 0, 0, 0, 1, 1, 0, 1, 0, 1, 0, 1, 0, 0, 0, 0, 1, 0, 0, 1,
1, 0, 0, 1, 1, 1, 0, 0, 1, 0, 0, 1, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0,
1, 0, 0, 0, 0, 1, 0, 1, 1, 0, 0, 0, 0, 1, 0, 1, 0, 1, 0, 1, 0,
0, 1, 0, 0, 1, 0, 0, 0, 0, 1, 1, 0, 1, 0, 1, 0, 0, 0, 0, 1, 1, 0,
0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0,
1, 0, 0, 0, 1, 0, 0, 0, 1, 1, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1,
0, 1, 1, 1, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 1, 0, 0, 1,
0, 0, 1, 0, 1, 1, 0, 0, 1, 0, 0, 1, 1, 0, 0, 1, 0, 0, 1, 0, 0,
0, 0, 0, 1, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 1, 0, 0, 1,
1, 1, 1, 0, 0, 1, 1, 1, 0, 1, 0, 1, 0, 1, 0, 0, 0, 0, 1, 0, 0],
dtype=int64)
```

```
In [12]: from sklearn.model_selection import StratifiedShuffleSplit as sss
sss = sss(test_size=0.25, random_state=0)
ss = sss.get_n_splits(X, y)
ss
```

Out[12]: 10

```
In [13]: for r,(train_index,test_index) in enumerate (sss.split(X,y)):
    print(r)
    print("Train    :",train_index)
    print("Test     :",test_index)

0
Train    : [432 453 706 606 118 421   3 157 400 497 296   30 201 314   32 475 501 293
681 381 346   34 119 642 248 397 101 645 498 659 182 154 531 210 437 528
763 533 666 382 160 197 176 555  18 699 518  29 585 404 486 332 650   33
 75 766 510 207 580 656 683 423 385 436 704 185 292 216 517 304 179 126
 97  39 608   31 167   85 256 647 424 581 653 470 617 260 224 741 250 193
217 204 233 543 625 302 509 371 223 440 549 365 552  37 258 646 500  82
384 262  44 758 348 511 383 760 430 689 394  28 120 412 635  17 220 150
724 406   2 364 698 445 226 722 367 135 271   86 718  73 428 513 275 759
707  13 350 621  23 629 597 558 249 672 599 756 410 477 491 172 715  40
184 420 447 206 578 251 560 143 557  57 316 326   61 244 730  95 173 211
106 246 512 257  52 640 317  22 564 561 419 343 214 312 196 155 130 213
141 750 283 618 468 527 269 186 754 255 731 502 396 716 376 664 146 753
401 170 748   83 110 733 209   7 630 408 624 417 116 728 550 124 690 321
612 366 631 131 341 200 586   70 450 127 572  12 614 767 102 493   4   1
446 516 494 615 335 700 290   79 322 278 158 474 422 569 411 337  19 306
402 115 121 616 601 660   65 177 161 547 610  14 591 347 149 602 194 472
488 693 388 229 300 389 559 342 462   93 105 476 478 655 454   64 499 584
523 729 764 311 709 327 266 222 480 589 461 323 637   35  41 238 174 372
544 730 606 334 305 637 670 606 735  66 433 640 111 761 753 370 100 670
```

```
In [14]: from sklearn.model_selection import train_test_split as tts
X_train,X_test,y_train,y_test=tts(X,y,test_size=.25,random_state=42)
```

```
In [15]: from sklearn.linear_model import LogisticRegression as lr
```

```
In [16]: logi=lr(penalty='l2',C=10.0)
```

```
In [17]: logi.fit(X_train,y_train)
```

C:\Users\user\anaconda3\lib\site-packages\sklearn\linear_model_logistic.py:814: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression (https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

n_iter_i = _check_optimize_result(

Out[17]: LogisticRegression(C=10.0)

```
In [18]: y_pred=logi.predict(X_test)
y_pred
```

Out[18]: array([0, 0, 0, 0, 0, 0, 0, 1, 1, 1, 0, 1, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0,
1, 0, 1, 1, 0, 0, 0, 0, 1, 1, 1, 1, 1, 1, 0, 1, 1, 0, 1, 1, 0,
0, 1, 1, 0, 0, 1, 0, 1, 1, 0, 0, 0, 1, 0, 0, 1, 1, 0, 0, 0, 1,
0, 1, 0, 1, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 1, 0,
0, 0, 0, 0, 0, 1, 1, 1, 0, 0, 1, 0, 1, 0, 1, 1, 1, 0, 0, 1, 0, 1,
0, 0, 0, 1, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 1, 1, 1,
0, 0, 1, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0,
0, 0, 0, 1, 1, 0, 1, 1, 0, 0, 0, 1, 0, 0, 1, 1, 1, 0, 0, 1, 1, 0,
0, 0, 0, 0, 1, 1, 0, 1, 1, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0], dtype=int64)

Step3. predict on a new sample

```
In [19]: pre=logi.predict([[6,200,90,10,25,23.3,0.672,42]])
if pre==0:
    print("Non-diabetic patient",pre)
else:
    print("Diabetic patient",pre)
```

Diabetic patient [1]

C:\Users\user\anaconda3\lib\site-packages\sklearn\base.py:450: UserWarning: X does not have valid feature names, but LogisticRegression was fitted with feature names
warnings.warn(

Compute Classification Metrics

```
In [20]: from sklearn.metrics import *
```

```
In [21]: accuracy_score(y_test, y_pred)
```

```
Out[21]: 0.7291666666666666
```

```
In [22]: precision_score(y_test,y_pred)
```

```
Out[22]: 0.6164383561643836
```

```
In [23]: recall_score(y_test,y_pred)
```

```
Out[23]: 0.6521739130434783
```

```
In [24]: roc_auc_score(y_test,y_pred)
```

```
Out[24]: 0.7122658183103571
```


Step4. Understand Correlation

```
In [25]: from sklearn.metrics import confusion_matrix as cfm
cf=cfmx(y_test,y_pred)
cf
```

```
Out[25]: array([[95, 28],
               [24, 45]], dtype=int64)
```

```
In [26]: import seaborn as sb
sb.heatmap(cf,annot=True)
```

```
Out[26]: <AxesSubplot:>
```



Step5. Normalization using MinMaxScaler and rebuild LoR

```
In [ ]:
```

```
In [27]: from sklearn.preprocessing import MinMaxScaler as mms
ms=mms()
ms_X_train=ms.fit_transform(X_train)
ms_X_train
```

```
Out[27]: array([[0.76470588, 0.64824121, 0.          , ..., 0.59463487, 0.20964987,
                0.38333333],
               [0.23529412, 0.64824121, 0.70491803, ..., 0.52309985, 0.06532878,
                0.03333333],
               [0.17647059, 0.30653266, 0.67213115, ..., 0.51266766, 0.0704526 ,
                0.41666667],
               ...,
               [0.58823529, 0.50753769, 0.70491803, ..., 0.67958271, 0.45175064,
                0.28333333],
               [0.          , 0.70854271, 0.          , ..., 0.6318927 , 0.05422716,
                0.13333333],
               [0.          , 0.6281407 , 0.78688525, ..., 0.33532042, 0.07856533,
                0.          ]])
```

```
In [28]: ms_X_test=ms.transform(X_test)
ms_X_test
```

```
Out[28]: array([[0.35294118, 0.49246231, 0.47540984, ..., 0.50670641, 0.15029889,
                0.36666667],
               [0.11764706, 0.56281407, 0.6147541 , ..., 0.53204173, 0.02988898,
                0.          ],
               [0.11764706, 0.54271357, 0.52459016, ..., 0.45901639, 0.03415884,
                0.          ],
               ...,
               [0.35294118, 0.61809045, 0.59016393, ..., 0.50074516, 0.27967549,
                0.21666667],
               [0.17647059, 0.3919598 , 0.40983607, ..., 0.46199702, 0.07258753,
                0.08333333],
               [0.17647059, 0.53266332, 0.59016393, ..., 0.38450075, 0.05508113,
                0.1          ]])
```

```
In [29]: ms_lr=lr()
ms_lr.fit(ms_X_train,y_train)
ms_y_pred=ms_lr.predict(ms_X_test)
print("Predictions of scaled data using MinMaxScaler:",ms_y_pred)
```

```
Predictions of scaled data using MinMaxScaler: [0 0 0 0 0 0 0 1 1 0 1 0 0 0 0 0 0 1 0 0 0 1 1 0 0 0 0 1 1 1 0 0 1 1
0 1 0 0 1 1 0 0 1 1 0 0 1 0 1 1 0 0 0 1 0 0 1 1 0 0 0 0 0 0 1 0 1 1 0 0 0
0 1 0 0 0 0 1 0 0 0 0 1 1 0 0 0 0 0 0 0 0 1 0 0 1 0 1 0 1 0 0 1 0 0 0
0 0 1 0 0 1 0 0 0 0 0 0 0 0 1 0 1 1 1 1 0 0 1 0 0 1 1 0 0 0 0 0 0 0
0 1 0 0 0 0 0 0 0 1 1 0 0 1 0 0 0 1 0 0 1 1 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 1 0 0 0]
```

```
In [30]: accuracy_score(y_test, ms_y_pred)
```

```
Out[30]: 0.7291666666666666
```

```
In [31]: precision_score(y_test,ms_y_pred)
```

```
Out[31]: 0.639344262295082
```

```
In [32]: recall_score(y_test,ms_y_pred)
```

```
Out[32]: 0.5652173913043478
```

```
In [33]: roc_auc_score(y_test,y_pred)
```

```
Out[33]: 0.7122658183103571
```

Step6. Normalization using StandardScaler and rebuild LoR

```
In [34]: from sklearn.preprocessing import StandardScaler
ss=StandardScaler()
ss_X_train=ss.fit_transform(X_train)
ss_X_train
```

```
Out[34]: array([[ 2.80346794,  0.25977903, -3.78077929, ...,  1.03974028,
                  0.29608546,  0.96352088],
               [ 0.07832678,  0.25977903,  0.89724451, ...,  0.40945373,
                -0.70087555, -0.86295593],
               [-0.22446668, -1.85825286,  0.67966201, ...,  0.31753694,
                -0.66548048,  1.13747105],
               ...,
               [ 1.89508755, -0.61235174,  0.89724451, ...,  1.78820556,
                  1.96850229,  0.44167036],
               [-1.13284707,  0.63354937, -3.78077929, ...,  1.36801453,
                -0.77756486, -0.34110542],
               [-1.13284707,  0.13518892,  1.44120077, ..., -1.24504846,
                -0.6094383 , -1.03690611]])
```

```
In [35]: ss_X_test=ss.transform(X_test)
ss_X_test
```

```
Out[35]: array([[ 0.6839137 , -0.70579433, -0.625833 , ...,  0.26501306,
                -0.11390738,  0.87654579],
                [-0.52726014, -0.26972894,  0.29889263, ...,  0.48823955,
                -0.94569142, -1.03690611],
                [-0.52726014, -0.39431905, -0.29945925, ..., -0.15517797,
                -0.91619553, -1.03690611],
                ...,
                [ 0.6839137 ,  0.07289387,  0.13570575, ...,  0.21248918,
                0.77981801,  0.09377001],
                [-0.22446668, -1.32874488, -1.06099801, ..., -0.12891603,
                -0.65073254, -0.60203068],
                [-0.22446668, -0.45661411,  0.13570575, ..., -0.81172646,
                -0.77166568, -0.51505559]])
```

```
In [36]: ss_lor=lr()
ss_lor.fit(ss_X_train,y_train)
ss_y_pred=ss_lor.predict(ss_X_test)
```

```
In [37]: ss_y_pred
```

```
Out[37]: array([0, 0, 0, 0, 0, 0, 0, 1, 1, 1, 0, 1, 0, 0, 0, 0, 1, 1, 0, 0,
                1, 0, 1, 1, 0, 0, 0, 1, 1, 1, 1, 1, 1, 0, 1, 1, 0, 1, 1, 0,
                0, 1, 1, 0, 0, 1, 0, 1, 1, 0, 0, 0, 1, 0, 0, 1, 1, 0, 0, 0, 0, 1,
                0, 1, 0, 1, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 1, 0,
                0, 0, 0, 0, 0, 1, 1, 1, 0, 0, 1, 0, 1, 0, 1, 1, 1, 0, 0, 1, 0, 1,
                0, 0, 0, 1, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 1, 1, 1, 1,
                0, 0, 1, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0,
                0, 0, 0, 1, 1, 0, 1, 1, 0, 0, 0, 1, 0, 0, 1, 1, 1, 0, 0, 1, 1, 0,
                0, 0, 0, 0, 1, 1, 0, 0, 0, 1, 0, 0, 1, 1, 1, 0, 0, 1, 1, 0,
                0, 0, 0, 0, 1, 1, 0, 1, 1, 0, 0, 0, 1, 0, 0, 0, 1, 1, 0,
                0, 0, 0, 0, 1, 1, 0, 1, 1, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0])
```

```
In [38]: accuracy_score(y_test,ss_y_pred)
```

```
Out[38]: 0.7291666666666666
```

```
In [39]: precision_score(y_test,ss_y_pred)
```

```
Out[39]: 0.6164383561643836
```

```
In [40]: recall_score(y_test,ss_y_pred)
```

```
Out[40]: 0.6521739130434783
```

```
In [41]: roc_auc_score(y_test,ss_y_pred)
```

```
Out[41]: 0.7122658183103571
```


Step7. Plot ROC curve

```
In [42]: pred_prob1=ss_lor.predict_proba(ms_X_test)
```

```
In [43]: fpr1, tpr1, thresh1 = roc_curve(y_test, pred_prob1[:,1], pos_label=1)
```

```
In [44]: import matplotlib.pyplot as plt
plt.plot(fpr1,tpr1,linestyle='-',color='blue',label='MinMaxScaler values')
```

```
Out[44]: [<matplotlib.lines.Line2D at 0x233af7234c0>]
```



Step8. Comparison with KNN classifier

```
In [45]: from sklearn.neighbors import KNeighborsClassifier as KNC
knn=KNC(n_neighbors=4)
knn=knn.fit(X_train,y_train)
```

```
In [46]: knn_y_pred=knn.predict(X_test)
```

```
In [47]: from sklearn.preprocessing import MinMaxScaler as ms
mi=ms()
mi_X_train=mi.fit_transform(X_train)
mi_X_train
```

```
Out[47]: array([[0.76470588, 0.64824121, 0.          , ..., 0.59463487, 0.20964987,
                0.38333333],
               [0.23529412, 0.64824121, 0.70491803, ..., 0.52309985, 0.06532878,
                0.03333333],
               [0.17647059, 0.30653266, 0.67213115, ..., 0.51266766, 0.0704526 ,
                0.41666667],
               ...,
               [0.58823529, 0.50753769, 0.70491803, ..., 0.67958271, 0.45175064,
                0.28333333],
               [0.          , 0.70854271, 0.          , ..., 0.6318927 , 0.05422716,
                0.13333333],
               [0.          , 0.6281407 , 0.78688525, ..., 0.33532042, 0.07856533,
                0.          ]])
```

```
In [48]: mi_X_test=mi.transform(X_test)
mi_X_test
```

```
Out[48]: array([[0.35294118, 0.49246231, 0.47540984, ..., 0.50670641, 0.15029889,
                0.36666667],
               [0.11764706, 0.56281407, 0.6147541 , ..., 0.53204173, 0.02988898,
                0.          ],
               [0.11764706, 0.54271357, 0.52459016, ..., 0.45901639, 0.03415884,
                0.          ],
               ...,
               [0.35294118, 0.61809045, 0.59016393, ..., 0.50074516, 0.27967549,
                0.21666667],
               [0.17647059, 0.3919598 , 0.40983607, ..., 0.46199702, 0.07258753,
                0.08333333],
               [0.17647059, 0.53266332, 0.59016393, ..., 0.38450075, 0.05508113,
                0.1          ]])
```

```
In [49]: mi_knn=KNC()
mi_knn=mi_knn.fit(mi_X_train,y_train)
```

```
In [50]: mi_y_pred=mi_knn.predict(mi_X_test)
mi_y_pred
```

```
Out[50]: array([0, 0, 0, 1, 1, 0, 0, 1, 1, 0, 0, 1, 0, 0, 0, 1, 0, 0, 1, 0, 1, 0,
                0, 0, 1, 0, 0, 0, 0, 0, 1, 1, 1, 1, 1, 1, 1, 0, 0, 1, 0, 1, 1, 0,
                0, 0, 1, 0, 0, 1, 0, 1, 1, 0, 0, 0, 0, 0, 1, 1, 1, 0, 0, 1, 0, 0,
                0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 1, 0, 0, 0, 1, 0, 0, 0, 0, 1, 1, 0,
                0, 0, 0, 0, 0, 1, 1, 1, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 1, 0, 1,
                0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 1, 1, 1,
                0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1,
                0, 1, 0, 1, 1, 0, 1, 0, 0, 0, 0, 1, 0, 1, 1, 1, 0, 1, 0, 1, 0,
                0, 1, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 1, 1, 0, 0])
```

```
In [51]: accuracy_score(y_test,mi_y_pred)
```

```
Out[51]: 0.6770833333333334
```

```
In [52]: precision_score(y_test,mi_y_pred)
```

```
Out[52]: 0.5522388059701493
```

```
In [53]: recall_score(y_test,mi_y_pred)
```

```
Out[53]: 0.5362318840579711
```

```
In [54]: roc_auc_score(y_test,mi_y_pred)
```

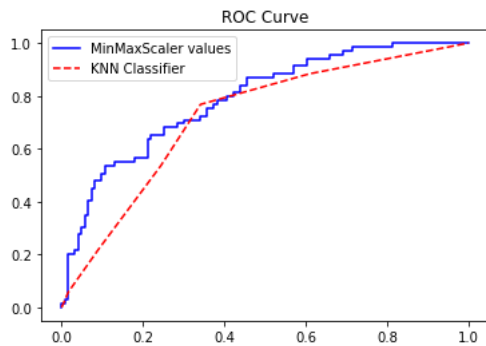
```
Out[54]: 0.6461647225167905
```

Step9. Update ROC curve

```
In [55]: pred_prob2=mi_knn.predict_proba(mi_X_test)

fpr2,tpr2,thresh2=roc_curve(y_test,pred_prob2[:,1],pos_label=1)
```

```
In [56]: plt.plot(fpr1,tpr1,linestyle='-',color='blue',label='MinMaxScaler values')
plt.plot(fpr2,tpr2,linestyle='--',color='red',label='KNN Classifier')
plt.title('ROC Curve')
plt.legend(loc='best')
plt.savefig('ROC',dpi=300)
plt.show()
```



Step10. Regularization

```
In [57]: from sklearn.linear_model import LogisticRegressionCV
model1=LogisticRegressionCV(Cs=10,cv=4,penalty='l1',solver='liblinear')
model2=LogisticRegressionCV(Cs=10,cv=4,penalty='l2')
```

```
In [58]: model1.fit(ms_X_train,y_train)
model2.fit(ms_X_train,y_train)
```

```
Out[58]: LogisticRegressionCV(cv=4)
```

```
In [59]: rg_y_pred1 = model1.predict(ms_X_test)
rg_y_pred2 = model2.predict(ms_X_test)
```

```
In [60]: from sklearn.metrics import roc_auc_score
l1_auc = roc_auc_score(y_test, rg_y_pred1)
l1_auc = (' LOR L1 MINMAX AUC', l1_auc)
l1_auc
```

```
Out[60]: (' LOR L1 MINMAX AUC', 0.694591728525981)
```

```
In [61]: from sklearn.metrics import roc_auc_score
l2_auc = roc_auc_score(y_test, rg_y_pred2)
l2_auc = (' LOR L2 MINMAX AUC', l2_auc)
l2_auc
```

```
Out[61]: (' LOR L2 MINMAX AUC', 0.7090844821491693)
```

Step11. Update ROC curve

```
In [62]: pred_prb7 = model1.predict_proba(ms_X_test)
pred_prb8 = model2.predict_proba(ms_X_test)
```

```
In [63]: fpr,tbr,threshold = roc_curve(y_test, pred_prob1[:,1],pos_label=1)
fpr1,tbr1,threshold1 = roc_curve(y_test, pred_prob2[:,1],pos_label=1)
fpr2,tbr2,threshold2= roc_curve(y_test, pred_prb7[:,1],pos_label=1)
fpr3,tbr3,threshold3 = roc_curve(y_test, pred_prb8[:,1],pos_label=1)
```



```
In [64]: plt.plot(fpr, tbr, linestyle='-', color='blue', label='LogisticRegression')
plt.plot(fpr1, tbr1, linestyle='--', color='orange', label='KNN')
plt.plot(fpr3, tbr3, linestyle='-', color='midnightblue', label='LoR 12')
plt.plot(fpr2, tbr2, linestyle='-', color='gray', label='LoR 11')
plt.title('Receiver Operating Characteristic')
plt.legend(loc = 'best')
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')
plt.show()
```

